

```
import pandas as pd
import seaborn as sns
```

Read The DataSet

```
df = pd.read_csv('Churn_Modelling.csv')
```

```
df.shape
```

```
(10000, 14)
```

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
      'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
      'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
df
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	1	15634602	Hargrave	619	France	Female
42						
1	2	15647311	Hill	608	Spain	Female
41						
2	3	15619304	Onio	502	France	Female
42						
3	4	15701354	Boni	699	France	Female
39						
4	5	15737888	Mitchell	850	Spain	Female
43						
...
...						
9995	9996	15606229	Obijiaku	771	France	Male
39						
9996	9997	15569892	Johnstone	516	France	Male
35						
9997	9998	15584532	Liu	709	France	Female
36						
9998	9999	15682355	Sabbatini	772	Germany	Male
42						
9999	10000	15628319	Walker	792	France	Female
28						
	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	

1	1	83807.86	1	0	1
2	8	159660.80	3	1	0
3	1	0.00	2	0	0
4	2	125510.82	1	1	1
...
9995	5	0.00	2	1	0
9996	10	57369.61	1	1	1
9997	7	0.00	1	0	1
9998	3	75075.31	2	1	0
9999	4	130142.79	1	1	0

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

df.head()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0

2	113931.57	1
3	93826.63	0
4	79084.10	0

Distinguish The Feature And Target Set

```
# input data/ Features Set
x = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
        'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
```

```
# output data/ Target Set
y = df['Exited']
```

x


	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...	
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	
9999	792	28	4	130142.79	1	1	

	IsActiveMember	EstimatedSalary
0	1	101348.88
1	1	112542.58
2	0	113931.57
3	0	93826.63
4	1	79084.10
...
9995	0	96270.64
9996	1	101699.77
9997	1	42085.58
9998	0	92888.52
9999	0	38190.78

[10000 rows x 8 columns]

```
sns.countplot(x=y)
```

```
<Axes: xlabel='Exited', ylabel='count'>
```



```
y.value_counts()

Exited
0    7963
1     2037
Name: count, dtype: int64
```

Normalize The Train And Test Data

```
# Normalize The Train And Test Data
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

x_scaled = scaler.fit_transform(x)

x_scaled
array([[ -0.32622142,  0.29351742, -1.04175968, ...,  0.64609167,
         0.97024255,  0.02188649],
       [-0.44003595,  0.19816383, -1.38753759, ..., -1.54776799,
         0.97024255,  0.21653375],
       [-1.53679418,  0.29351742,  1.03290776, ...,  0.64609167,
```

```

-1.03067011,  0.2406869 ],
...,
[ 0.60498839, -0.27860412,  0.68712986, ..., -1.54776799,
  0.97024255, -1.00864308],
[ 1.25683526,  0.29351742, -0.69598177, ...,  0.64609167,
 -1.03067011, -0.12523071],
[ 1.46377078, -1.04143285, -0.35020386, ...,  0.64609167,
 -1.03067011, -1.07636976]])

```

Cross Validation - Train Test Data

```

# Cross Validation - Train Test Data
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x_scaled, y,
                                                    random_state=0, test_size=0.25)

x.shape
(10000, 8)

x_test.shape
(2500, 8)

x_train.shape
(7500, 8)

```

Initialize and build the model

```

# Initialize and build the model
from sklearn.neural_network import MLPClassifier

ann = MLPClassifier(hidden_layer_sizes=(100,100,100), random_state=0,
max_iter=100, activation='relu')

ann.fit(x_train, y_train)

C:\Users\Shubham\anaconda3\Lib\site-packages\sklearn\neural_network\
_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (100) reached and the optimization
hasn't converged yet.
  warnings.warn(

MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100,
random_state=0)

```

```
y_pred = ann.predict(x_test)
y_pred
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)


from sklearn.metrics import ConfusionMatrixDisplay,
classification_report, accuracy_score

y_test.value_counts()
Exited
0    1991
1     509
Name: count, dtype: int64
```

Print Accuracy Score And Confusion Matrices

```
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x1e15a2fc6e0>
```



```
accuracy_score(y_test, y_pred)
0.8228
```

imbalanced Classifier is this

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.92	0.89	1991
1	0.59	0.45	0.51	509
accuracy			0.82	2500
macro avg	0.73	0.68	0.70	2500
weighted avg	0.81	0.82	0.81	2500

to balace the dataset

```
!pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in c:\users\shubham\
anaconda3\lib\site-packages (0.12.3)
Requirement already satisfied: numpy>=1.17.3 in c:\users\shubham\
anaconda3\lib\site-packages (from imbalanced-learn) (1.26.4)
Requirement already satisfied: scipy>=1.5.0 in c:\users\shubham\
anaconda3\lib\site-packages (from imbalanced-learn) (1.13.1)
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\
shubham\anaconda3\lib\site-packages (from imbalanced-learn) (1.4.2)
Requirement already satisfied: joblib>=1.1.1 in c:\users\shubham\
anaconda3\lib\site-packages (from imbalanced-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\
shubham\anaconda3\lib\site-packages (from imbalanced-learn) (2.2.0)
```

```
from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state = 0)
```

```
x_res, y_res = ros.fit_resample(x,y)
```

```
y_res.value_counts()
```

```
Exited
```

```
1    7963
```

```
0    7963
```

```
Name: count, dtype: int64
```

This above data is balanced using Over Sampler techniques

```
# Normalize
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

x_scaled = scaler.fit_transform(x_res)

x_scaled
array([[ -0.29877723,  0.08418894, -1.01840607, ...,  0.6512495 ,
         1.08223556,  0.00817382],
       [ -0.4103938 , -0.01032629, -1.36135608, ..., -1.53550983,
         1.08223556,  0.20261687],
       [ -1.48597169,  0.08418894,  1.03929402, ...,  0.6512495 ,
        -0.92401325,  0.22674468],
       ...,
       [ -0.84671313,  1.02934128,  0.01044398, ...,  0.6512495 ,
        -0.92401325,  1.28878188],
       [ -0.96847667,  0.65128034, -0.67545605, ..., -1.53550983,
         1.08223556, -1.21851316],
       [ -1.5874413 ,  0.74579558,  1.03929402, ...,  0.6512495 ,
         1.08223556,  1.42417217]])

# Cross Validation - Train Test Data
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x_scaled, y_res,
                                                    random_state=0, test_size=0.25)

x_res.shape
(15926, 8)

x_test.shape
(3982, 8)

x_train.shape
(11944, 8)

# Initialize and build the model
from sklearn.neural_network import MLPClassifier

ann = MLPClassifier(hidden_layer_sizes=(100,100,100), random_state=0,
                    max_iter=100, activation='relu')

ann.fit(x_train, y_train)
```



```

C:\Users\Shubham\anaconda3\Lib\site-packages\sklearn\normalization\
_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic
Optimizer: Maximum iterations (100) reached and the optimization
hasn't converged yet.
  warnings.warn(

MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100,
random_state=0)

y_pred = ann.predict(x_test)
y_pred
array([1, 1, 1, ..., 1, 1, 1], dtype=int64)

from sklearn.metrics import ConfusionMatrixDisplay,
classification_report, accuracy_score


y_test.value_counts()

Exited
1    2001
0    1981
Name: count, dtype: int64

ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x1e15a48b260>

```



```
accuracy_score(y_test, y_pred)
```

```
0.865143144148669
```

balanced Classifier is this

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.90	0.82	0.86	1981
1	0.84	0.91	0.87	2001
accuracy			0.87	3982
macro avg	0.87	0.86	0.86	3982
weighted avg	0.87	0.87	0.86	3982